



GITGOOD
Help gitGood, gitBetter

DEPLOYMENT DIAGRAM DESCRIPTION

DEMO 3

WEATHER TO WEAR

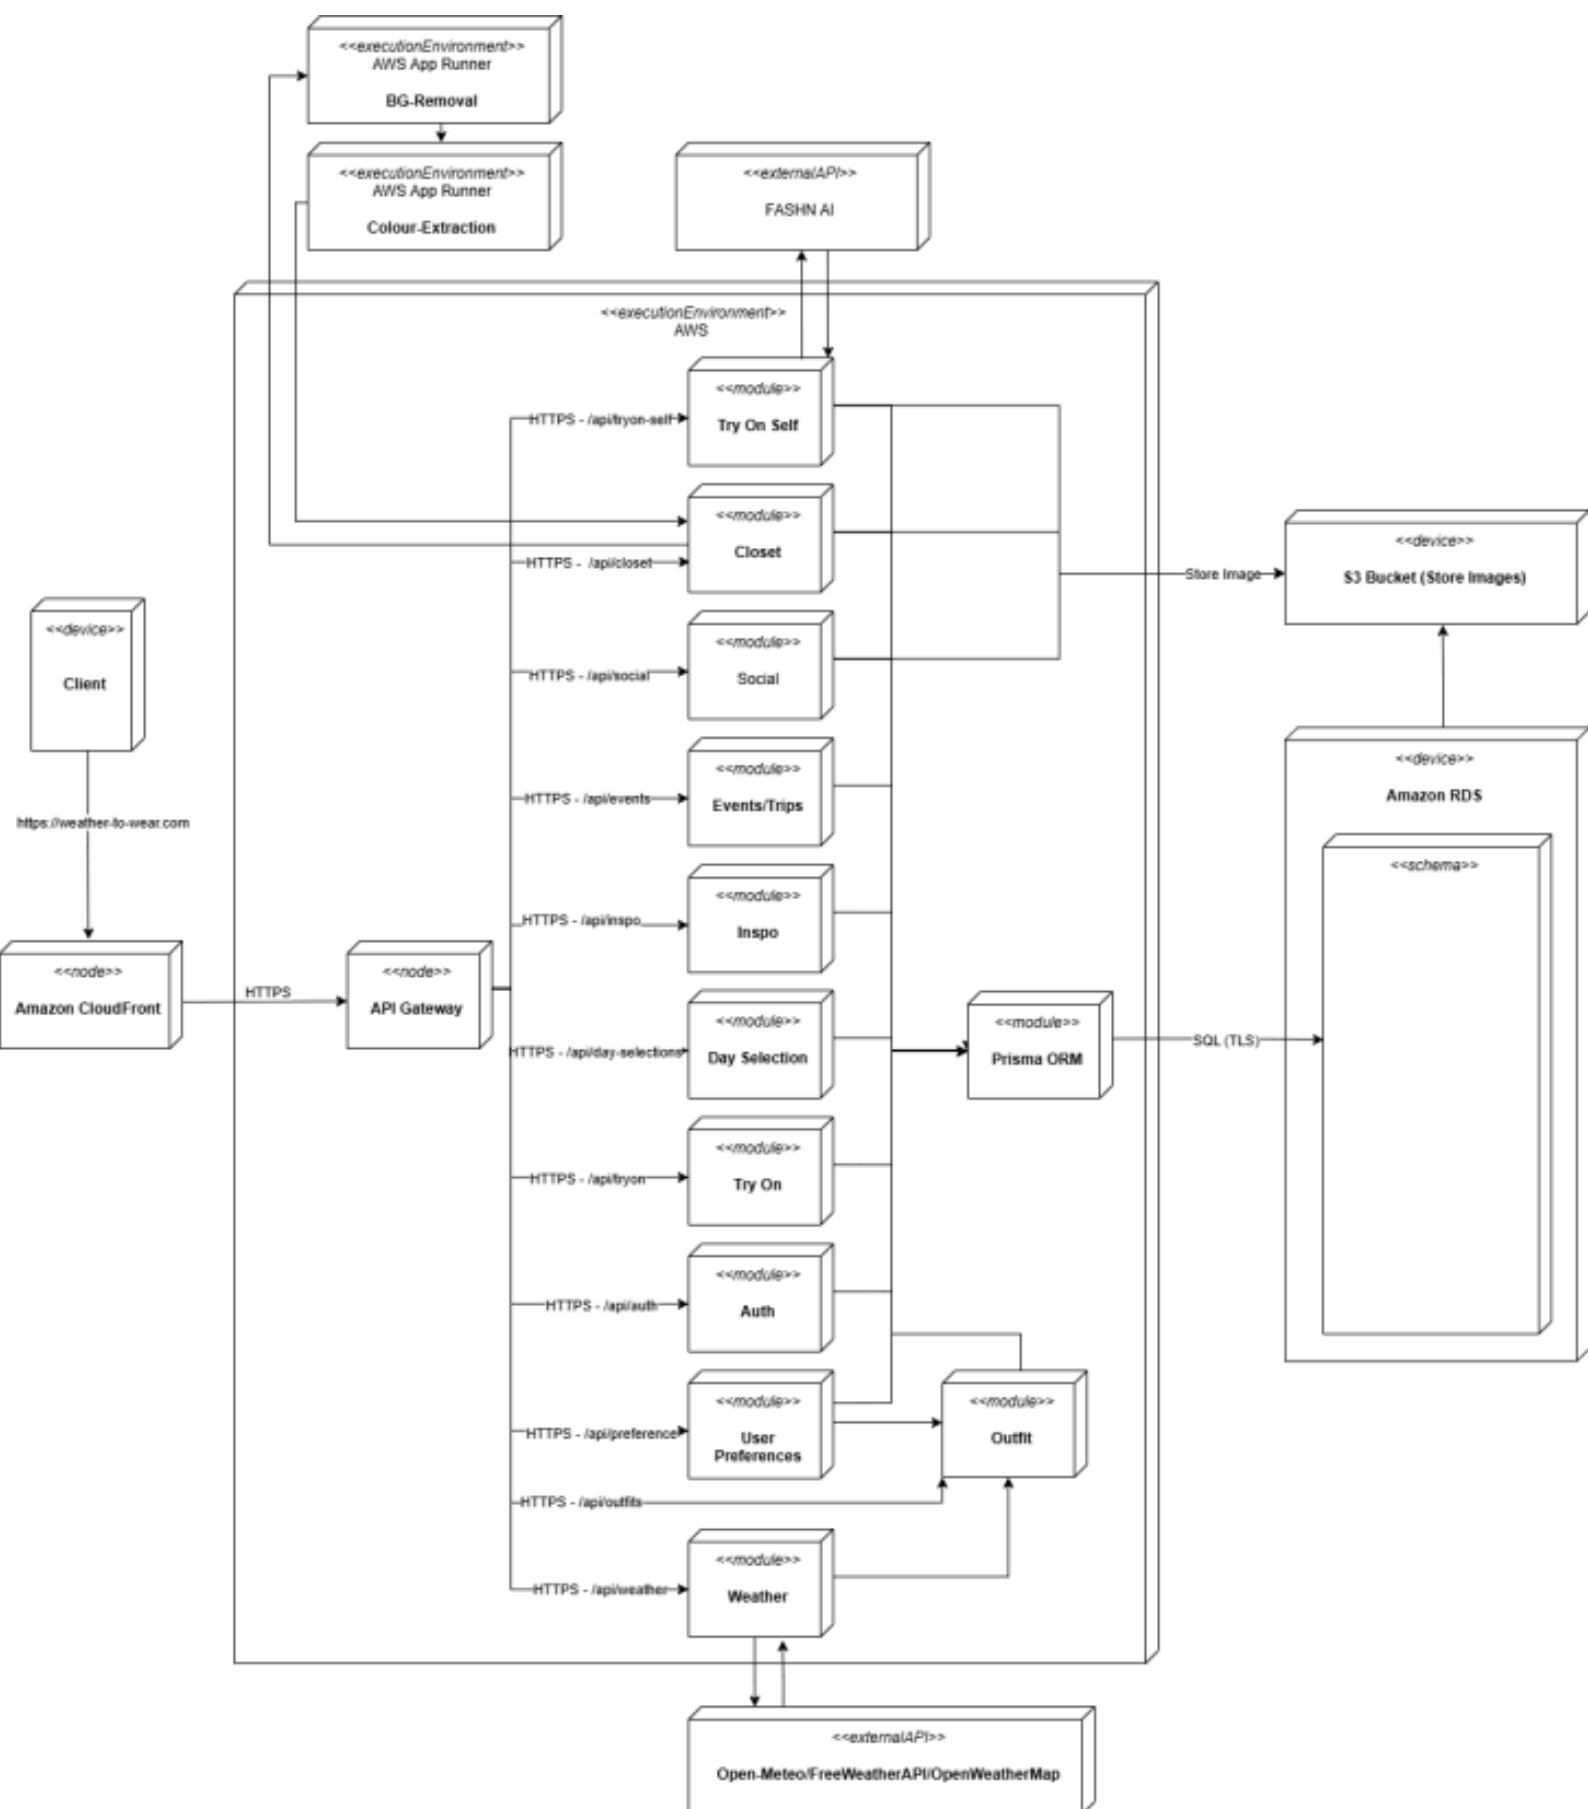


KYLE LIEBENBERG
DIYA BUDHIA
ALISHA PERUMAL

IBRAHIM SAID
TAYLOR SERGEL

Table of Contents

DIAGRAM	3
NODES, COMPONENTS, AND ARTIFACTS	4
RUNTIME COMMUNICATION PATHS	5
SECURITY CONTROLS & TRUST BOUNDARIES	6
AVAILABILITY, SCALABILITY, AND PERFORMANCE	6
OBSERVABILITY & OPERATIONS	7
CI/CD TOUCHPOINTS	7
SECURITY CONTROLS & TRUST BOUNDARIES	7
AVAILABILITY, SCALABILITY, AND PERFORMANCE	8





Legend

- `<<node>>` – execution environment or managed service (e.g., App Runner, CloudFront, RDS).
- `<<component>>` – deployable software unit (e.g., “Weather to Wear API”).
- `<<artifact>>` – binary/config/data item (e.g., Docker image, React build, secrets, user images).
- `<<interface>>` – interaction contract (REST/HTTPS endpoints).

Nodes, Components, and Artifacts

CLIENT EDGE

- Client `<<node>>`
- Browser/PWA that requests the frontend over HTTPS and calls the backend REST API.

FRONTEND DELIVERY

- Amazon CloudFront – Frontend CDN `<<node>>`
 - HTTPS `<<interface>>` to the client.
 - Serves the React PWA from the S3 origin using Origin Access Control (OAC).
- Amazon S3 – Frontend Bucket `<<node>>`
 - Build React App (build/) `<<artifact>>` uploaded by CI/CD.
 - Public access blocked; only CloudFront reads via OAC.
- GitHub Actions (OIDC) – Deploy Frontend `<<node>>`
 - Build Artifact `<<artifact>>` (React build output) synced to the frontend S3 bucket; CloudFront invalidation completes the release.

BACKEND COMPUTE

- AWS App Runner – Backend API `<<node>>`
 - Weather to Wear API (Node.js + Express + Prisma) `<<component>>`
 - REST API `v1/api/*` `<<interface>>` exposed to the client over HTTPS.
 - Docker image `weather-backend:prod` `<<artifact>>` pulled from ECR.
 - Reads secrets at start-up; runs `prisma migrate deploy` before serving traffic (ensures schema parity).
 - Environment (plain): `PORT`, `NODE_ENV`, `S3_BUCKET_NAME`, `S3_REGION`, `UPLOADS_CDN_DOMAIN`, `BG_REMOVAL_URL`, `COLOR_EXTRACT_URL`.

IMAGE PROCESSING MICROSERVICES

- AWS App Runner – bg-removal (public) <<node>>
 - Background Removal Service (U²-Net) <<component>>
 - HTTP POST /remove-bg <<interface>>
 - Docker image bg-removal:prod <<artifact>>.
- AWS App Runner – color-extract (public) <<node>>
 - Color Extraction Service (KMeans) <<component>>
 - HTTP POST /extract-colors <<interface>>
 - Docker image color-extract:prod <<artifact>>.

CONTAINER REGISTRY & SECRETS

- Amazon ECR <<node>>
 - OCI Images (tags: prod) <<artifact>> for backend and both microservices.
- AWS Secrets Manager <<node>>
 - Secrets <<artifact>>: DATABASE_URL, JWT_SECRET, and weather API keys.
 - Backend fetches via GetSecretValue at boot.

PRIVATE NETWORK & DATA STORES

- VPC: w2w-prod-vpc <<node>>
 - S3 Gateway Endpoint <<node>> for private S3 egress (saves NAT cost/latency).
 - App Runner VPC Connector <<node>> provides private connectivity to RDS on port 5432/TLS.
 - Private Subnets <<node>> host:
 - Amazon RDS PostgreSQL <<node>> (instance w2w-postgres-prod, port 5432, no public endpoint).
 - RDS Automated Backups <<artifact>>.
- Amazon CloudFront – Uploads CDN <<node>>
 - HTTPS <<interface>> for serving user images globally (via OAC to S3).
- Amazon S3 – Uploads Bucket <<node>>
 - User Images (PNG/JPEG/WebP) <<artifact>>, private with Block Public Access ON.
 - Backend writes via IAM; CloudFront reads via OAC.

Runtime Communication Paths

1. PWA delivery

- a. Client → CloudFront (Frontend) <<interface:HTTPS>> → S3 Frontend Bucket via OAC.
- b. The PWA is entirely static; API base URL is injected at build time.

2. User API traffic

- a. Client → App Runner – Backend API <<interface:REST /api/v1>>.
- b. Requests pass through the Application Layer (routing/validation/auth) and into domain services.

3. Secrets retrieval

- a. Backend API → Secrets Manager (GetSecretValue for DATABASE_URL, JWT_SECRET) at boot.

4. Database access

- a. Backend API → App Runner VPC Connector → RDS PostgreSQL on 5432/TLS (private).
- b. Prisma performs queries via Repositories in the Persistent Layer.

5. Media write path (uploads)

- a. Backend API (after processing) → S3 Uploads Bucket (Put/Delete via IAM).
- b. Stored object keys are returned as CDN URLs to the client.

6. Media read path (images)

- a. Client → CloudFront (Uploads) <<interface:HTTPS>> → S3 Uploads via OAC.
- b. Buckets remain private; no public ACLs/policies.

7. Image processing pipeline

- a. Backend API → bg-removal POST /remove-bg → color-extract POST /extract-colors → return metadata/bytes → S3 write → persist DB record.

8. Container images

- a. App Runner services pull Docker images <<artifact>> from ECR during deployment.

9. Migrations on start

- a. Backend container runs prisma migrate deploy at boot against RDS to ensure the live schema matches code (e.g., isTrip column).

Security Controls & Trust Boundaries

- **Transport security:** TLS/HTTPS on all external edges (CloudFront and App Runner).
- **Private data plane:** RDS in private subnets; reachable only via App Runner VPC Connector ENIs.
- **Least-privilege IAM:**
 - Backend role: read specific Secrets; write to the uploads bucket path; pull from ECR.
 - Microservices: only what they require (no DB access).
- **Private S3:** Both buckets have Block Public Access ON; CloudFront OAC is the only read path.
- **AuthN/AuthZ:** JWT for protected routes; RBAC for admin operations.
- **CORS:** allow-list the frontend origin only.
- **Config secrets:** stored in Secrets Manager; non-sensitive config via environment variables.

Availability, Scalability, and Performance

- **Stateless services** (backend & microservices) on **App Runner** support horizontal scale and **rolling deployments** with health checks.
- **CDN** terminates global image traffic, reducing API load and user-perceived latency.
- **RDS** automated backups; Multi-AZ can be enabled when budget permits.
- **Caches** (in-process, bounded TTL) reduce dependency on external weather providers and speed up repeated reads.

Observability & Operations

- App Runner logs for build/deploy/runtime; application logs to stdout/stderr.
- Optional CloudFront access logs to S3.
- Database monitoring via RDS Performance Insights/metrics.
- Alarms on latency/5xx/error budget; synthetic checks canaries during deploys.

CI/CD Touchpoints

- Frontend pipeline (GitHub Actions): build React → sync to S3 Frontend → CloudFront invalidation.
- Backend & microservices: build images → push to ECR → App Runner StartDeployment (or auto-deploy on new image tag).
- Migrations run on container start to keep RDS schema in sync.

Assumptions & Constraints

- RDS is not publicly accessible.
- All S3 buckets are private; access is via CloudFront OAC (reads) and backend IAM (writes).
- Microservices are public endpoints but callable only by the backend (documented; consider token/IP allow-list if needed).
- No message queue yet; the pipeline is synchronous today but the boundary allows adding SQS later without changing public APIs.

Failure Modes & Degradation Paths

- **Image service outage:**
 - backend returns a clear error or degrades to storing the original image; user can retry.
- **Weather provider failure:**
 - backend uses a fallback provider and/or cached summaries.
- **RDS transient errors:**
 - short retries with jitter on read-only paths; write paths remain idempotent to prevent duplication.
- **CDN miss or S3 latency:**
 - served on origin fetch; object then cached at the edge.